

Scientific visualization with ParaView

Part 2

Jarno van der Kolk
jvanderk@uottawa.ca



**Digital Research
Alliance** of Canada



- ✓ slides, data, codes at <https://jarno.ca/pv.zip>
 - the link will download a file pv.zip (~35MB)
 - unpack it to find codes/, data/ and slides{1,2}.pdf
 - command line: wget <https://jarno.ca/pv.zip> -O pv.zip
- ✓ install ParaView 5.13.x on your laptop from <http://www.paraview.org/download>

EXPORTING SCENES (PRE-COMPUTED POLYGONS)

ParaView Glance

<https://kitware.github.io/paraview-glance>

PV Glance is an open-source **standalone** web app for **in-browser 3D sci-vis**

- very easy to use, ideal for sharing pre-built 3D scenes via the web
 - no server ⇒ up to medium-size data (server support planned in future versions)
 - interactive manipulation of pre-computed polygons
 - volumetric images, molecular structures, geometric objects, point clouds
 - written in JavaScript and vtk.js + can be further customized with vtk.js and ParaViewWeb for custom web and desktop apps
 - source and installation instructions <https://github.com/kitware/paraview-glance>
-

1. Create a visualization with several layers, make **all layers visible in the pipeline**
2. Many options in **File** → **Export Scene...** ⇒ save as VTKJS to your laptop
3. Open <https://kitware.github.io/paraview-glance/app>
4. Drag the newly saved file to the dropzone on the website
5. Interact with individual layers in 3D: **rotate and zoom, change visibility, representation, variable, colourmap, opacity**

Automatically load a visualization into Glance

<https://discourse.paraview.org/t/customise-pv-glance/2831>

- Use the query syntax `GLANCEAPPURL?name=FILENAME&url=FILEURL` to pass `name` and `url` to the web server
- E.g. using ParaView Glance website <https://kitware.github.io/paraview-glance/app?name=sineEnvelope.vtkjs&url=https://raw.githubusercontent.com/razoumov/publish/master/data/sineEnvelope.vtkjs>
 - shortened to <https://bit.ly/2KtPWNf>
- You can parse long strings with JavaScript (next slide)

Embed your vis into a website with an iframe (`embed.html`)

```
<!DOCTYPE html>
<html>
  <head>
    <title>Sine envelope function</title>
  </head>
  <body>
    <h1>3D sine envelope function</h1>

    <script>
      var app = "https://kitware.github.io/paraview-glance/app";
      var dir = "https://raw.githubusercontent.com/razoumov/publish/master/data/";
      var file = "sineEnvelope.vtkjs";
      document.write("<iframe src='" + app + "?name=" + file + "&url=" +
                    dir + file +
                    "' id='iframe' width='1100' height='900'></iframe>");
    </script>

    <p>More stuff in here</p>
  </body>
</html>
```

- JavaScript here only to parse long strings

ANIMATION IN PARAVIEW

Animation methods

1. Use ParaView's built-in animation of any property of any pipeline object

- easily create snazzy animations, somewhat limited in what you can do
- in Time Manager: select object, select property, create a new track with “+”, double-click the track to edit it, press “▶”

Animation methods

1. Use ParaView's built-in animation of any property of any pipeline object

- easily create snazzy animations, somewhat limited in what you can do
- in Time Manager: select object, select property, create a new track with “+”, double-click the track to edit it, press “▶”

2. Use ParaView's ability to recognize a sequence of similar files

- time animation only, very convenient
- try loading data/2d*.vtk sequence and animating it (visualize one frame and then press “▶”)

Animation methods

1. Use ParaView's built-in animation of any property of any pipeline object

- easily create snazzy animations, somewhat limited in what you can do
- in Time Manager: select object, select property, create a new track with “+”, double-click the track to edit it, press “▶”

2. Use ParaView's ability to recognize a sequence of similar files

- time animation only, very convenient
- try loading `data/2d*.vtk` sequence and animating it (visualize one frame and then press “▶”)

3. Script your animation in Python (covered in next section)

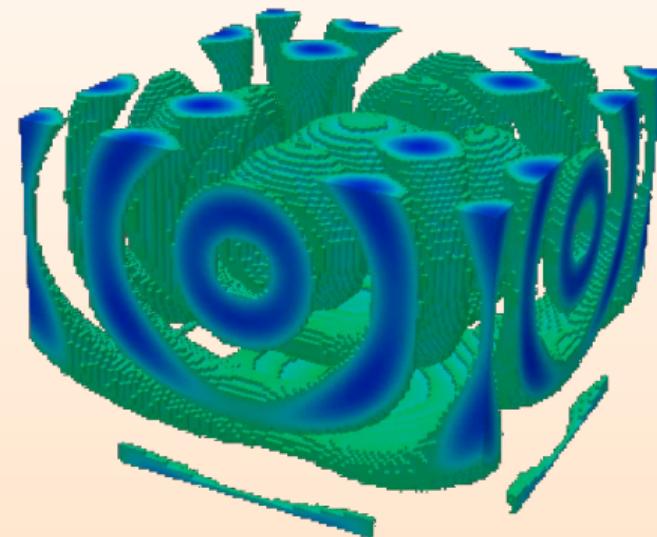
- steep learning curve, very powerful, can do anything you can do in the GUI
- typical usage scenario: generate one frame per input file
- a simpler exercise without input files: see next slide

Exercise: animating function growth

- 3D sine envelope wave function defined inside a unit cube ($x_i \in [0, 1]$)

$$f(x_1, x_2, x_3) = \sum_{i=1}^2 \left[\frac{\sin^2 \left(\sqrt{\xi_{i+1}^2 + \xi_i^2} \right) - 0.5}{\left[0.001(\xi_{i+1}^2 + \xi_i^2) + 1 \right]^2} + 0.5 \right], \text{ where } \xi_i \equiv 15(x_i - 0.5)$$

- Reproduce the movie using sineEnvelope.nc



Exercise: animating function growth (cont.)

To visualize a single frame of the movie:

1. load data/sineEnvelope.nc (discretized on a 100^3 grid)
2. apply Threshold keeping only data from 1.2 to 2
3. apply Clip: origin $O = (49.5, 15, 49.5)$, normal $N = (0, -1, 0)$
4. colour by the right quantity

Two possible solutions:

1. bring up **Time Manager** to animate Clip's O_2 from 0 to 99, for best results save animation as a sequence of PNG files (NOTE: bug in Paraview, need to set range manually...)
2. covered in the next section: Start/Stop Trace to record the workflow, save the corresponding **Python script**, enclose **parts of it** into a loop changing O_2 from 0 to 99 and writing a series of PNG screenshots, run it inside ParaView to produce 100 frames

in either case, merge PNGs using a 3rd-party tool, e.g.

```
ffmpeg -r 30 -i frame%04d.png -c:v libx264 -pix_fmt yuv420p \
-vf "scale=trunc(iw/2)*2:trunc(ih/2)*2" movie.mp4
```

Camera animation in the GUI

Good introductory resource https://www.paraview.org/Wiki/Advanced_Animations

1. Start with any static visualization
2. Click on 'Adjust Camera' icon (one of the left-side icons on top of the visualization window)
 - adjust / write down Camera Focal Point
3. Bring up Time Manager (or erase all previous timelines)

(3a) In Time Manager:

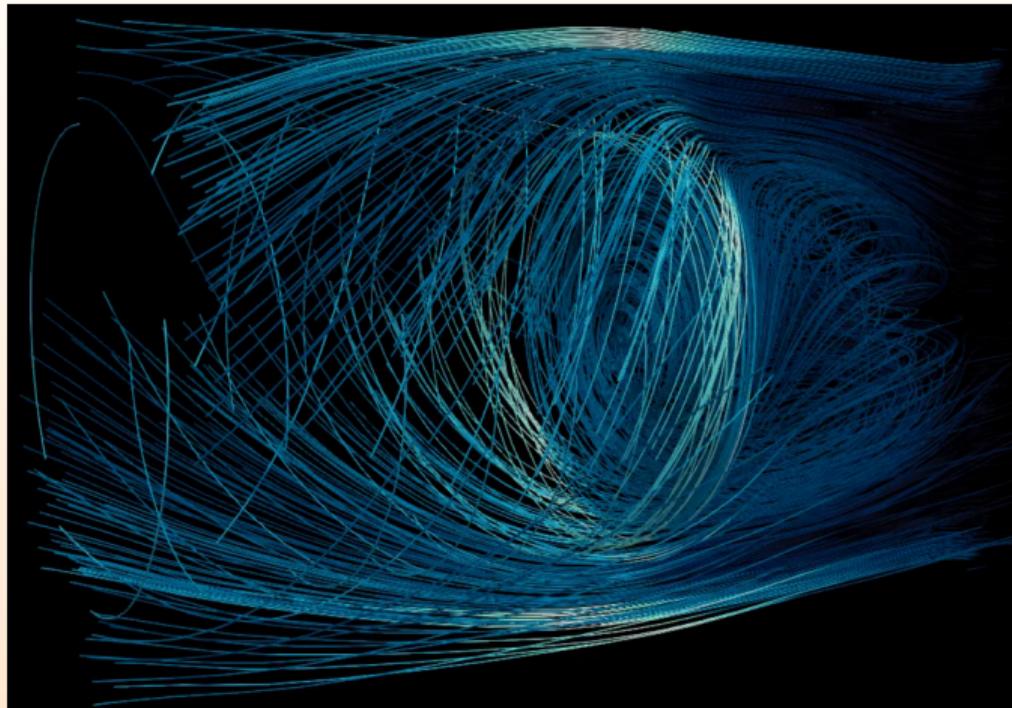
- select Camera - Follow Path
- click "+" to create a new timeline
- set Center = Camera Focal Point, for the rest accept default settings
- adjust the number of frames

(3b) In Time Manager:

- select Camera - Follow Path
- click "+" to create a new timeline
- double-click on the white (or black) timeline
- double-click on Path... in the right column
- click on Camera Position
 - a yellow path with spheres will appear
 - drag the spheres around
- also can change Camera Focus and Up Direction

4. Click "▶"

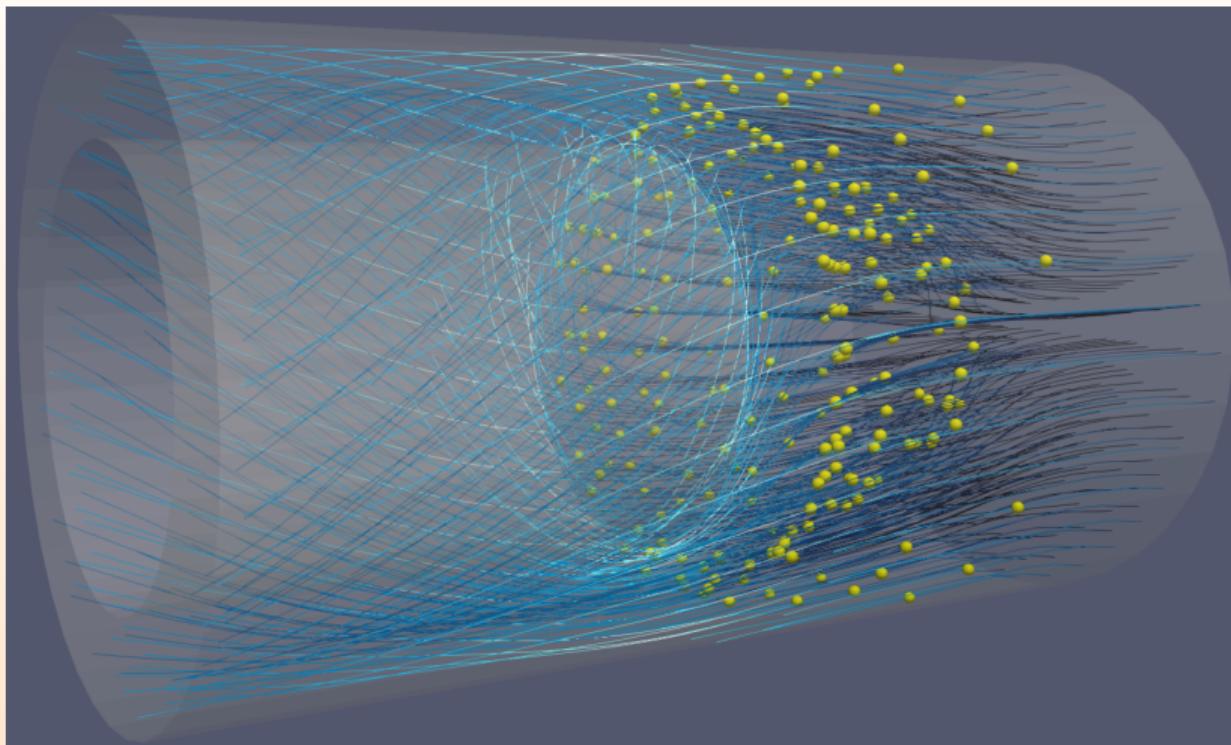
Animating stationary flow: streamlines through a slice



Animating stationary flow: streamlines through a slice (cont.)

1. Load `disk_out_ref.ex2` making sure to load velocity
2. Draw a radius-z plane slice through the center, origin $O = (0, 0, 0)$ and normal $N = (0, 0, 1)$
3. Stream Tracer With Custom Source: `input=disk_out_ref.ex2, seedSource=Slice1`
4. Tube filter with $r = 0.02$
5. Time Manager: animate Slice's O_2 from 0 to 2 (full range [-5.75,5.75])
6. Use 100 frames, black background, blue2cyan colourmap, colour with vorticity
7. Unselect "Show Plane"
8. Save animation as PNGs, encode at 10 fps

Animating a stationary flow: time contours

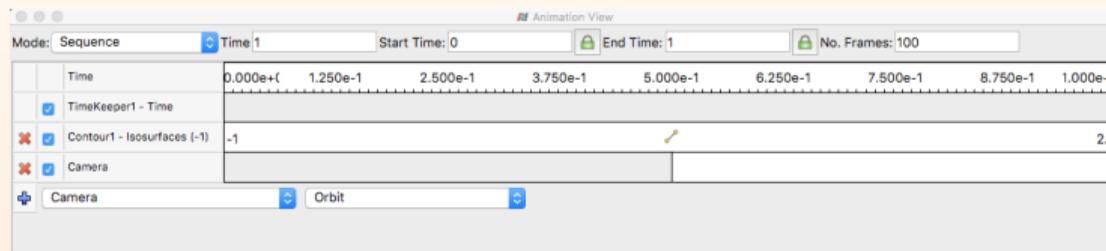


Animating a stationary flow: time contours (cont.)

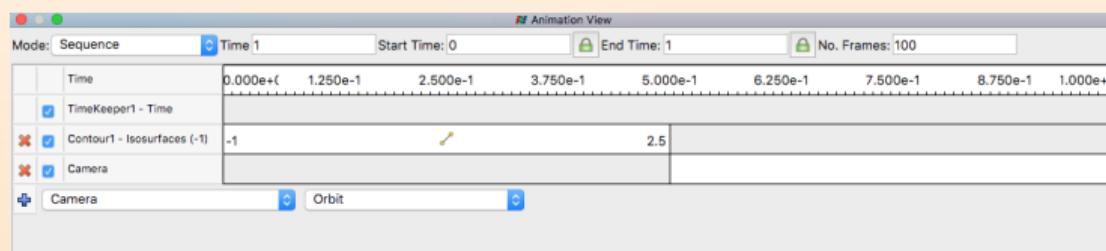
1. Start with the streamtracer lines, however drawn
2. Apply a Contour filter to the output of Streamtracer
 - contour by Integration Time
 - probe the range of values that works best
3. Apply Glyph filter to the output of Contour
4. Time Manager: animate Contour | Isosurfaces
5. This video was recorded with 2000 frames at 60 fps
 - such high resolution only for the final production video
 - debugging animation with 100 frames is perfectly Ok

Exercise: several timelines in one animation

1. Start with the previous integration-time-contour animation
2. Add the second timeline to the animation: Camera - Follow Path from $t = 0.5$ to $t = 1$ (while the first animation is still playing for its second half)

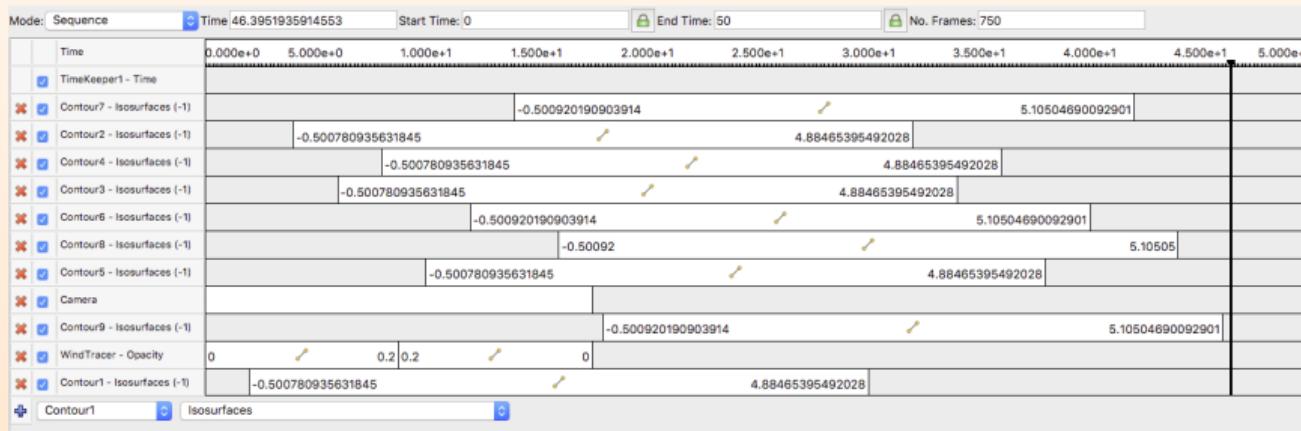


3. Now complete integration-time-contour animation before rotation



Combining many timelines in one animation (cont.)

- In principle, can add as many timelines (with their individual time intervals and variables!) to the animation as you want
- Here is an example from WestGrid's 2017 *Visualize This* competition submission by Nadya Moisseeva (UBC)



PYTHON SCRIPTING IN PARAVIEW

Batch scripting for automating visualization

Official documentation at https://www.paraview.org/Wiki/ParaView/Python_Scripting

- Why use scripting?

- automate mundane or repetitive tasks, e.g., making frames for a movie
- document and store your workflow
- use ParaView on clusters from the command line and/or via batch jobs

- In the GUI: View | Python Shell opens a Python interpreter

- write or paste your script there
- use the button to run an external script from a file

- `[/usr/bin/ /usr/local/bin/ /Applications/Paraview*.app/Contents/bin/]`

pvython will give you a Python shell connected to a ParaView server (local or remote) without the GUI

- `[/usr/bin/ /usr/local/bin/ /Applications/Paraview*.app/Contents/bin/]`

pvybatch --force-offscreen-rendering `script.py` is a serial (on some machines parallel) application using a local ParaView server  **make sure to save your visualization**

- `[/usr/bin/ /usr/local/bin/ /Applications/Paraview*.app/Contents/MacOS/]`

`paraview --script=codes/displayWireframe.py` to start ParaView GUI and auto-run the script

First script

- Bring up View | Python Shell
- “Run Script” codes/displaySphere.py

displaySphere.py

```
from paraview.simple import *

sphere = Sphere() # create a sphere pipeline object

print(sphere.ThetaResolution) # print one of the attributes of the sphere
sphere.ThetaResolution = 16

Show() # turn on visibility of the object in the view
Render()
```

- Can always get help from the command line

```
help(paraview.simple)      # will display a help page on paraview.simple module
help(Sphere)
help(Show)
help(sphere)    # to see this object's attributes
dir(paraview.simple)
```

Using filters

- “Run Script” codes/displayWireframe.py

displayWireframe.py

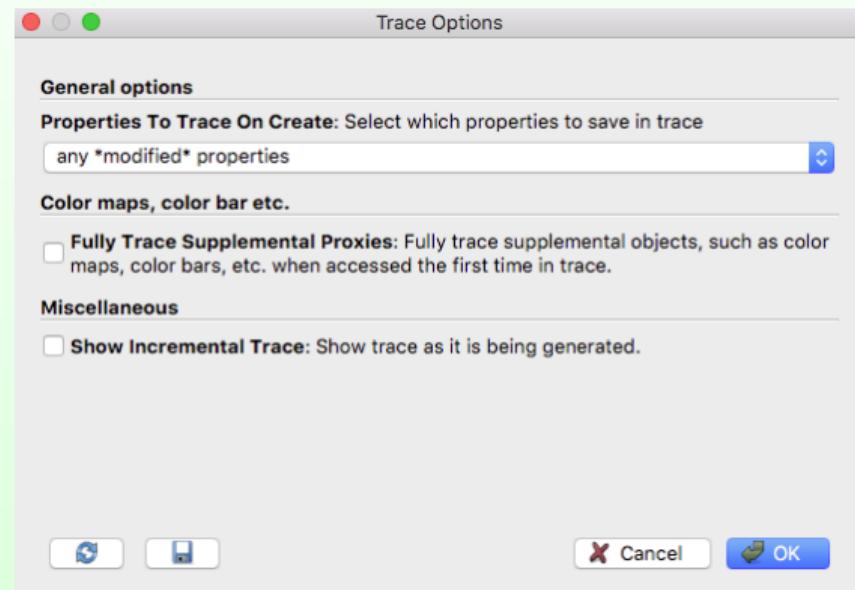
```
from paraview.simple import *
sphere = Sphere(ThetaResolution=36, PhiResolution=18)
wireframe = ExtractEdges(Input=sphere) # apply Extract Edges to sphere
Show() # turn on visibility of the last object in the view
Render()
```

- Try replacing Show() with Show(sphere)
- Also try replacing Render() with SaveScreenshot('/path/to/wireframe.png') and running via pvbatch

Trace tool

Generate Python code from GUI operations

- Newer ParaView: Tools | Start / Stop Trace
- Older ParaView: Tools | Python Shell | Trace | Start / Stop / Show Trace



Passing information down the pipeline

... and other useful high-level workflow functions

- `GetSources()` gets a list of pipeline objects
- `GetActiveSource()` gets the active object
- `SetActiveSource()` sets the active object
- `GetRepresentation()` returns the *view representation* for the active pipeline object and the active view
- `GetActiveCamera()` returns the active camera for the active view
- `GetActiveView()` returns the active view
- `CreateRenderView()` creates standard 3D render view
- `ResetCamera()` resets the camera to include the entire scene but preserve orientation (or does nothing ☺)

There is quite a bit of overlap between these two:

```
help(GetActiveCamera())
help(GetActiveView())
```

Camera animation with scripting

1. Let's load data/sineEnvelope.nc and draw an isosurface at $\rho = 0.15$
2. Compare the focal point to the center of rotation (must be the same for object to stay in view)

```
v1 = GetActiveView()  
print(v1.CameraFocalPoint)  
print(v1.CenterOfRotation)  
  
if not => ResetCamera()
```

3. Look up azimuthal rotation

```
dir(GetActiveCamera())  
help(GetActiveCamera().Azimuth)
```

4. Rotate by 10° around the view-up vector

```
camera = GetActiveCamera()  
camera.Azimuth(10)  
Render()
```

Camera animation: full rotation

✍ Can paste longer commands from clipboard.txt

5. Do full rotation and save to disk

```
nframes = 360
for i in range(nframes):
    print(v1.CameraPosition)
    camera.Azimuth(360./nframes)      # rotate by 1 degree
    SaveScreenshot('/path/to/frame%04d'%(i)+'.png')
```

6. Merge all frames into a movie at 30 fps

```
ffmpeg -r 30 -i frame%04d.png -c:v libx264 -pix_fmt yuv420p \
-vf "scale=trunc(iw/2)*2:trunc(ih/2)*2" spin.mp4
```

Camera animation: flying towards the focal point

1. Optionally reset the view manually or with `ResetCamera()`
2. Now let's fly 2/3 of the way towards the focal point

```
initialCameraPosition = v1.CameraPosition[:]    # force a real copy
nframes = 100
for i in range(nframes):
    coef = float(i+0.5)/float(1.5*nframes)    # runs from 0 to 2/3
    print(coef, v1.CameraPosition)
    v1.CameraPosition = [((1.-coef)*a + coef*b) \
        for a, b in zip(initialCameraPosition,v1.CameraFocalPoint)]
    SaveScreenshot('/path/to/out%04d'%(i)+'.png')
```

3. Create a movie

```
ffmpeg -r 30 -i out%04d.png -c:v libx264 -pix_fmt yuv420p \
    -vf "scale=trunc(iw/2)*2:trunc(ih/2)*2" approach.mp4
```

Exercise: write and run a complete off-screen script

1. Mac/Linux/Windows: create a script with standalone ParaView GUI

- use Start/Stop Trace
- load `data/sineEnvelope.nc` and draw an isosurface at $\rho = 0.15$
- save the image as PNG

2. Test-run your script with `pvbatch` on your laptop

```
$ pvbatch --force-offscreen-rendering script.py
```

- **Linux:** `pvbatch` should be in one of your system's bin directories
- **Mac:** `pvbatch` should be in `/Applications/ParaView*.app/Contents/bin`
- **Windows:** `pvbatch` does not exist (or so I am told), but you can use `pypython`
 - ⇒ you will need to locate it yourself
- You can also run this script on one of Magic Castle with

```
$ module load paraview/5.13.1
$ pvbatch --force-offscreen-rendering script.py
```

3. Modify the script to create some animation

Extracting data from VTK objects

Do this from *View | Python Shell* or from *pvcpython* (either shell will work)

```
# codes/extractValues.py
from paraview.simple import *

dir = '/Users/razoumov/training/paraviewWorkshop/data/'
data = NetCDFReader(FileName=[dir+'stvol.nc'])
local = servermanager.Fetch(data) # get the data from the server
print(local.GetNumberOfPoints())

for i in range(10):
    print(local.GetPoint(i)) # coordinates of first 10 points

pd = local.GetPointData()
print(pd.GetArrayName(0)) # the name of the first array

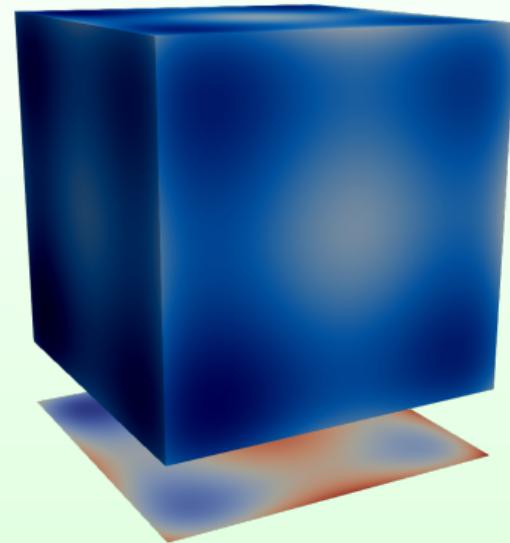
result = pd.GetArray('f(x,y,z)')
print(result.GetDataSize())
print(result.GetRange())

for i in range(10):
    print(result.GetValue(i)) # values at first 10 points
```

This is useful for post-processing, e.g., feeding these into **numpy arrays** and doing further calculations in a Python script

Creating/modifying VTK objects

Let's say we want to plot a projection of a cubic dataset like `sineEnvelope.nc` along one of its principal axes, or do some other transformation for which there is no filter



- Calculator / Python Calculator filter cannot modify the geometry ...

Programmable filter

Watch our webinar <https://bit.ly/programmablefilter>

1. Apply Programmable Filter with OutputDataSetType = vtkUnstructuredGrid
2. Paste the following code codes/projectionUnstructured.py into the filter
(this code was last tested in ParaView 5.12.1)

```
numPoints = inputs[0].GetNumberOfPoints()
side = int(round(numPoints**(.1/3.)))      # round() in this Python returns float type
layer = side*side
rho = inputs[0].PointData['density']        # 1D flat array
points = vtk.vtkPoints()                  # create vtkPoints instance, to contain 100^2 points in the projection
proj = vtk.vtkDoubleArray(); proj.SetName('projection')    # create the projection array
for i in range(layer):                  # loop through 100x100 points
    x, y = inputs[0].GetPoint(i)[0:2]
    z, column = -20., 0.
    for j in range(side):
        column += rho.GetValue(i+layer*j)
    points.InsertNextPoint(x,y,z)          # also points.InsertPoint(i,x,y,z)
    proj.InsertNextValue(column)          # add value to this point

output.SetPoints(points)                  # add points to vtkUnstructuredGrid
output.GetPointData().SetScalars(proj)    # add projection array to these points

quad = vtk.vtkQuad()                    # create a cell
output.Allocate(side, side)             # allocate space for side^2 'cells'
for i in range(side-1):
    for j in range(side-1):
        quad.GetPointIds().SetId(0,i+j*side)
        quad.GetPointIds().SetId(1,(i+1)+j*side)
        quad.GetPointIds().SetId(2,(i+1)+(j+1)*side)
```

Using 3rd-party libraries from ParaView's Python

- `pypython` includes few common 3rd-party libraries such as `numpy`, `scipy`, `pandas`
- What if you want to use other libraries that were not bundled with ParaView?

Using 3rd-party libraries from ParaView's Python

- `pvpython` includes few common 3rd-party libraries such as `numpy`, `scipy`, `pandas`
- What if you want to use other libraries that were not bundled with ParaView?

1. Let's assume you work on a CC cluster; check your ParaView's Python version

```
module load StdEnv/2023 paraview/5.13.1
pvpython    # let's assume it says Python 3.11.5
```

2. Load the closest Python module, create a virtual env. and install your library there

```
module avail python    # python/3.11.5 is one of them
module load python/3.11.5
virtualenv --no-download astro    # this will install a new virtual environment into ~/astro
source ~/astro/bin/activate
pip install --no-index --upgrade pip
pip install --no-index xarray    # install an external package into this new environment
```

3. Next time you log in to the cluster, start `pvpython`:

```
module load StdEnv/2023 paraview/5.13.1
pvpython
```

4. Load your new virtual environment directly from Python:

```
filename = '/home/username/astro/bin/activate_this.py'
exec(open(filename).read(), {'__file__': filename})
from paraview.simple import *
import xarray    # this xarray comes from your new virtual environment
```

5. For batch workflows, replace `pvpython` with `pvbatch`

REMOTE AND DISTRIBUTED VISUALIZATION

Visualizing remote data

If your dataset is on a remote cluster, there are several options:

✗ download data to your desktop and visualize it locally
- limited by the dataset size and your desktop's CPU/GPU + memory

✗ run ParaView remotely on a larger machine via X11 forwarding
- your desktop $\xrightarrow{\text{ssh -Y}}$ larger machine running ParaView
- remote OpenGL apps will run either (1) software rasterizer on the cluster (usually the default) or (2) on your laptop's GPU (need to re-enable INdirect GLX inside X11 server and set `LIBGL_ALWAYS_INDIRECT=1`)

✓ run ParaView remotely on a larger machine via remote desktop

- your desktop $\xrightarrow{\text{VNC}}$ larger machine running ParaView
- you can always start a VNC server on an interactive cluster compute node by hand as described in our documentation <https://bit.ly/startVNC>
- remote OpenGL apps will run either (1) using software rasterizer on the cluster (usually the default) or (2) on cluster's GPU(s) via VirtualGL wrapper (see our VNC docs)
- the VNC slide is coming up

✓ run ParaView in client-server mode

ParaView client on your desktop \Rightarrow ParaView server on larger machine

✓ run ParaView via a GUI-less batch script (interactively or scheduled)

- render server can run with GPU rendering or purely in software
- data/render servers can run on single-core, or across several cores/nodes with MPI
- for interactive GUI work on clusters you should schedule interactive jobs, as opposed to running on the login nodes

Special remote vis cases

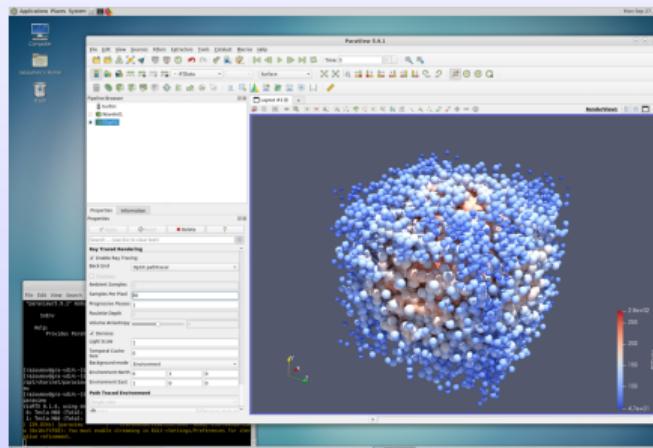
1. **In-situ visualization** = instrumenting a simulation code on the cluster to
 - 1.1 output graphics and/or
 - 1.2 act as on-the-fly server for a visualization frontend (ParaView/VisIt client on your laptop)
 - need to use a special library (ParaView's Catalyst or VisIt's libsim)
 - very advanced topic for another time
2. **Web-based visualization** with data served from another location

ParaView via remote desktop

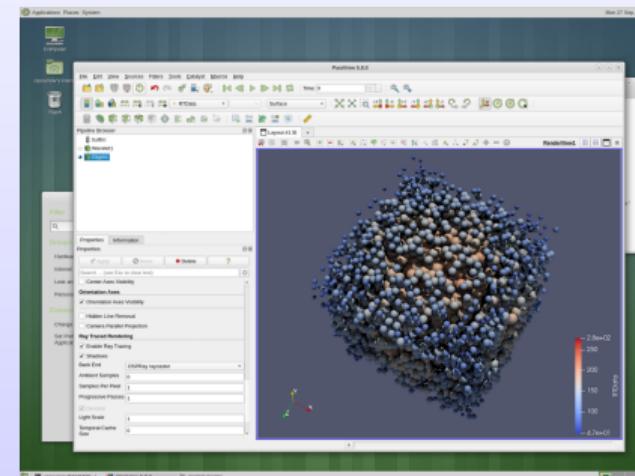
<https://docs.alliancecan.ca/wiki/VNC> or <https://docs.alliancecan.ca/wiki/JupyterHub>

You have several options:

- (1) run a VNC server on compute nodes with SSH tunnelling, connect via a VNC client
https://docs.alliancecan.ca/wiki/VNC#Compute_Nodes



- (2) VNC on `gra-vdi.computeCanada.ca`, connect via a VNC client



- (3) Remote Desktop via JupyterHub on Béluga, point your web browser at <https://jupyterhub.beluga.computeCanada.ca>

Cedar, Graham, Béluga, Narval clusters

- General-purpose CC clusters for a variety of workloads
 - entered production in phases since June 2017
 - located at SFU, UofWaterloo, École de technologie supérieure (Montreal)
 - 101,568 / 44,444 / 39,120 / 80,720 CPUs
 - many hundred NVIDIA GPUs (with 12GB/16GB/32GB on-board memory)
 - multiple types of nodes, with 128GB/256GB/0.5TB/1.5TB/3TB memory
 - specs at <https://docs.alliancecan.ca/wiki/Cedar>
(replace Cedar with Graham or Beluga or Narval)
- Batch-oriented environment for parallel and serial jobs ⇒ use Slurm scheduler and workload manager
- Identical software setup https://docs.alliancecan.ca/wiki/Available_software

Interactive jobs on Cedar / Graham / Béluga

- Client-server workflow is by definition interactive
- On Cedar interactive jobs should automatically go to one of Slurm's interactive partitions (CPU or GPU)

```
$ sinfo -p cpubase_interac
    # will list nodes and their states (idle, mixed, allocated, ...)
```

- `salloc` without a script name will start an interactive shell inside a submitted job on a compute node

```
$ salloc --time=1:0:0 --ntasks=4 ... --account=def-someuser
$ echo $SLURM_...      # access Slurm variables
$ module load ...      # set your environment
$ ./serialCode
$ srun ./mpiCode       # run an MPI code
$ exit                 # terminate the job (go back to the login node)
```

- You might need to specify `pvserver --server-port=11112` (etc.) if someone else is already using the default port 11111 on the same node

Question 1: should I use CPUs or GPUs for rendering?

- Can render on GPUs (*hardware acceleration*) or CPUs (*software rendering*) with both interactive and batch visualizations
 - GPUs have traditionally been faster for rendering graphics
 - in recent years better open-source software rendering libraries such as OSPRay (Intel's ray tracing) and OpenSWR (Intel's rasterizer) have largely closed the performance gap for many types of visualizations
- ⇒ I recommend starting with CPU rendering since you already likely have many CPUs! (see next slide)
- One might have to resort to software rendering if no GPUs are available, e.g., all taken by GP-GPU jobs
- I suggest doing **all hands-on exercises with CPU rendering**; also included slides on **GPU rendering** on the cluster

Question 2: how many CPUs/GPUs do I need?

- How many processors do we need? From *ParaView documentation*:
 - structured data (Structured Points, Rectilinear Grid, Structured Grid): one CPU core per \sim 20 million cells
 - unstructured data (Unstructured Points, Polygonal Data, Unstructured Grid): one CPU core per \sim 1 million cells
- Your main bottlenecks will be **physical memory** and **disk read speed**, and to a lesser extent **CPU/GPU rendering time** \Rightarrow to simplify things, to decide on the number of CPU cores for initial dataset exploration, use the dataset size
 - consider 80 GB dataset
 - base nodes have 128 GB memory with 32 cores \Rightarrow 3.5 GB/core (accounting for the OS, system tools, etc.) \Rightarrow 23 cores for this dataset
 - need to account for filters (and other processing), MPI buffers \Rightarrow minimum 32 cores
 - for comfortable processing with complex filters use 48 – 64 cores
- On large HPC systems ParaView is known to scale to $\sim 10^{12}$ cells (Structured Points) on \sim 10,000 cores and beyond
- Always do a scaling study before attempting to visualize large datasets
- It is important to understand **memory requirements of filters**
 - a typical structured \rightarrow unstructured filter increases memory footprint by \sim 3X

Remote Render Threshold

In ParaView's preferences can set (Render View | Remote/Parallel Rendering Options | Remote Render Threshold) beyond which rendering will be remote

- **default 20MB** ⇒ small rendering will be done on your laptop's GPU, interactive rotation with a mouse will be fast, but anything modestly intensive (under 20MB) will be shipped to your laptop and might be slow
- **0MB** ⇒ all rendering (including rotation) will be remote, so you will be really using the cluster's CPU(s)/GPU(s) for everything
 - good for large data processing
 - not so good for interactivity, especially on a slower connection
- experiment with the threshold to find a suitable value

Next few pages: remote rendering exercises

Short version:

1. create your visualization via interactive client-server using CPU rendering
2. save your visualization to PNG

Long version:

1. create your visualization via interactive client-server using CPU rendering
2. save your visualization to PNG
3. convert this workflow into a Python script
4. upload this Python script to the cluster
5. try running the script inside an interactive (`salloc`) job; debug if needed
6. once happy with the result, write a Slurm job submission script and submit this rendering as a batch (`sbatch`) job

Exercise 1 (on Graham): deep impact dataset

Dataset from IEEE 2018 SciVis Contest

- Dataset from *Deep Water Impact* simulation by John Patchett (LANL) and Galen Gisler (Univ. of Oslo)

- dataset details at <https://bit.ly/2Sxmjsq>
 - you can work with 269 low-resolution ($460 \times 280 \times 240$) snapshots in time
 - the original simulation is much higher resolution

- You can render this dataset in serial

- try to adapt the client-server instructions from “Parallel software rendering” slide (forward a few pages) to render on **one CPU**

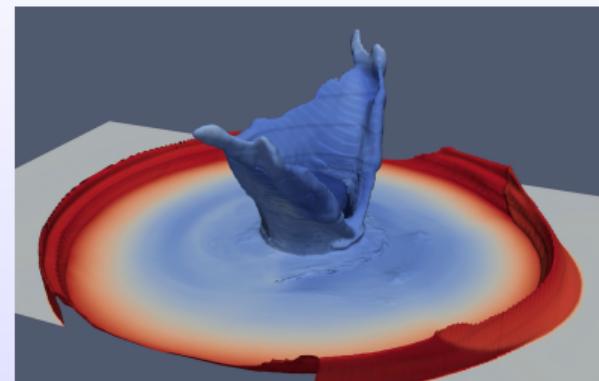
- Data in graham:/project/6024405/jarno/ieeevis2018/460x280x240 (115GB uncompressed in total)

- to simplify navigating to the dataset in ParaView, I highly recommend creating a symbolic link:

```
[cedar]$ mkdir -p ~/data
[cedar]$ ln -s /project/6024405/jarno/ieeevis2018/460x280x240 ~/data/deepImpact
```

- Demo on our training cluster

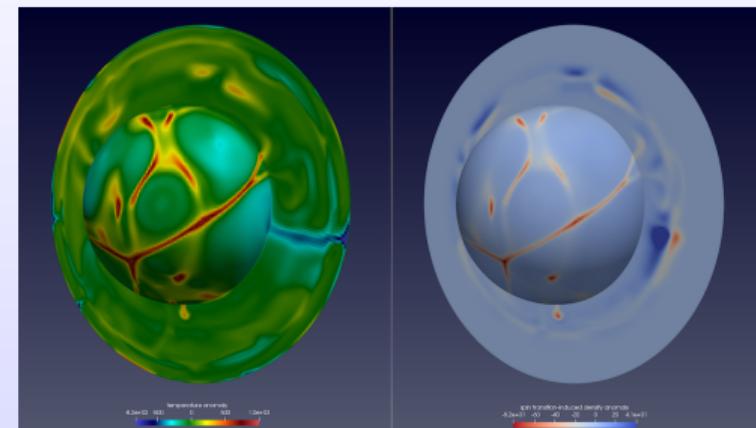
- compressed dataset (every 10th file in the timeline) \Rightarrow 12GB
 - demo on 8 cores on the training cluster
 - load all 26 frames, Contour by v02 (water volume fraction), colour by ρ , Rescale to data range



Exercise 2 (on Graham): Earth's mantle convection

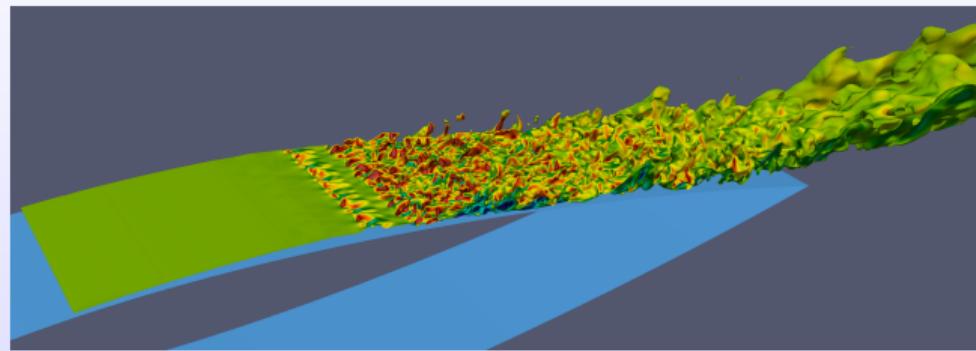
Dataset from IEEE 2021 SciVis Contest <https://scivis2021.netlify.app>

- Dataset from *Earth's Mantle Convection* simulation by Hosein Shahnas and Russell Pysklywec (U. of Toronto)
 - dataset details at <https://scivis2021.netlify.app/data>
 - 251 timesteps on a spherical $180 \times 201 \times 360$ grid
- You can render this dataset in serial
 - try to adapt the client-server instructions from “Parallel software rendering” slide (forward a few pages) to render on **one CPU**
- Data in graham:/project/6024405/jarno/ieeevis2021/spherical (89GB in total)
- Create a symbolic link to simplify navigating to the dataset in ParaView



Exercise 3 (on Graham): airflow over a turbine blade

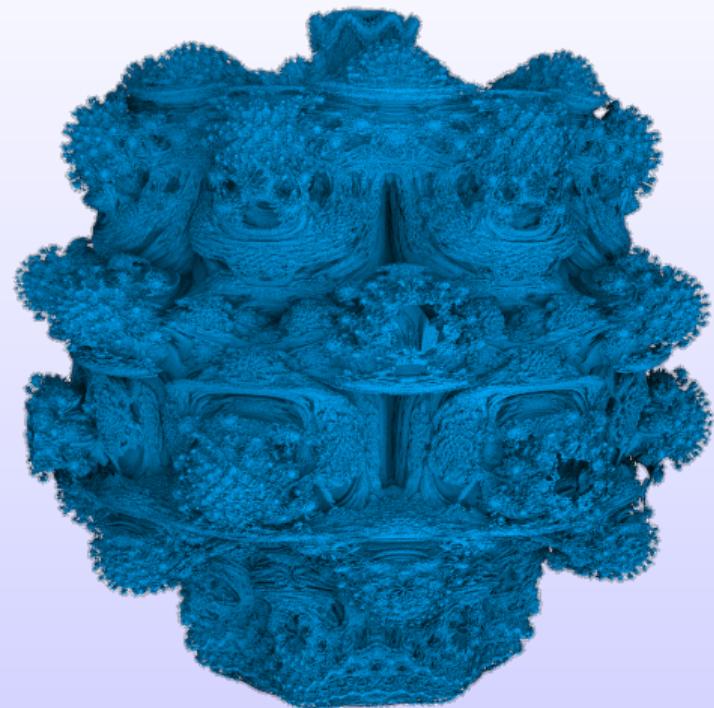
Dataset from WestGrid's 2019 <https://compute-canada.github.io/visualizeThis>



- OpenFOAM *decomposed* dataset: 512 cores, 86 timesteps, 5 hydro variables, ~1TB in total
 - kindly provided for this competition by Joshua Brinkerhoff (UBC Okanagan)
 - unstructured mesh ⇒ loading a single timestep from the **3D internal mesh** requires 200GB+ physical RAM
 - the **2D airfoil mesh** takes only 13.7 GB virtual memory for 1 timestep + 1 variable
 - data in graham:/project/6024405/jarno/visThis2019
- Image at the top shows the isosurface of constant air speed coloured by the Y-component of the vorticity, full animation rendering (86 timesteps) took 17 minutes on 128 Cedar CPU cores
- Create a symbolic link to simplify navigating to the dataset in ParaView

Exercise 4 (on the training cluster): Mandelbulb

- Visualize power-8 **Mandelbulb**
- Use the file `mandelbulb800.nc` – now sampled at 800^3
- Use 4–8 CPU cores on the training cluster via `salloc`
 1. consult the next three pages, use critical thinking – you will need to modify some of the commands!
 2. try to recreate the picture on the right: pay attention to the **lights** and **shadows**
 3. use `View → Memory Inspector` to keep an eye on memory usage
 4. optionally colour your dataset by `processID`



```
$ unzip /home/razoumov/shared/paraview.zip data/mandelbulb800.nc
$ ls -lh data/mandelbulb800.nc
```

Parallel software rendering

From interactive client-server debugging to remote batch rendering

1. On the cluster start remote parallel ParaView server:

```
$ cd scratch    # necessary on Graham
$ module load StdEnv/2023 paraview/5.11
$ salloc --time=0:60:0 --ntasks=128 --mem-per-cpu=3600 --account=def-someuser
$ mpirun -np 128 pvserver
```

2. Wait for it to start waiting for incoming connection:

```
Waiting for client...
Connection URL: cs://gra288:11111
Accepting connection(s): gra288:11111
```

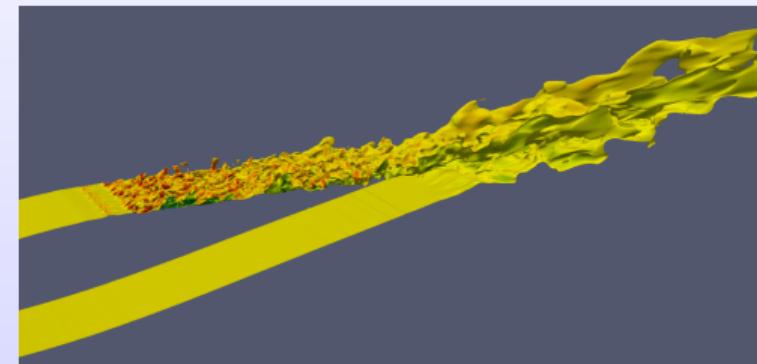
3. On your laptop start SSH port forwarding:

```
$ ssh graham.alliancecan.ca -L 11111:gra288:11111  # use the actual compute node
```

4. On your laptop start ParaView 5.13.x, click Connect, then connect to
`cs://localhost:11111`

Parallel software rendering (cont.)

5. **Tools** → **Start Trace**
6. Load OpenFOAM data, set Case Type = Decomposed
7. Apply Calculator: speed = mag(U)
8. Apply Contour at speed=0.8
9. Colour by (vorticity)_y
10. Load *Rainbow Desaturated* colourmap
11. Save the image as a PNG file
12. **Tools** → **Stop Trace**



13. Save the generated script as `airflow.py` locally
 - edit it in a text editor, simplify (most generated lines will be setting defaults)
 - provide the correct output PNG path on the remote system

Parallel software rendering (cont.)

14. Upload the script to the cluster:

```
$ scp airflow.py graham.alliancecan.ca:scratch/
```

15. On the cluster try running it as a parallel interactive job:

```
$ cd ~/scratch
$ salloc --time=0:60:0 --ntasks=128 --mem-per-cpu=3600 --account=def-someuser
$ module load gcc/9.3.0 paraview/5.13.1
$ mpirun -np 128 pvbatch --force-offscreen-rendering airflow.py
```

16. Once you are happy with the result, write a Slurm job submission script and submit it with sbatch

OpenGL context for off-screen rendering on a GPU

To render on a GPU from an OpenGL application such as ParaView, **traditionally you would require:**

1. OpenGL support in the GPU driver, and
2. an X server that handles windows and surfaces onto which client APIs can draw
 - run X11 server (typically started by root) on the GPU compute node, set DISPLAY=:0.\$gpuindex (get GPU index from Slurm)

Latest NVIDIA GPU drivers include EGL (*Embedded-System Graphics Library*) support enabling creation of an OpenGL context for off-screen rendering without an X server.

- Your OpenGL application needs to be **recompiled with EGL support** ⇒ use a special version of ParaView for GPU rendering without an X server; currently compiled into a module `paraview/5.13.1` that provides both **pvserver** for client-server and **pvbatch** for batch rendering
- Unlike X11, EGL does not require any special setting to scale to very high resolutions, e.g., 4K (3840×2160) – simply ask it to render a 4K image

Interactive client-server rendering on a cluster's GPU

Details in <http://bit.ly/2wrSvKV>

1. On Cedar/Graham/Béluga **submit an interactive job** to the GPU partition, e.g., a serial job:

```
$ salloc --time=0:30:0 --ntasks=1 --gpus-per-node=[type:]count \
--mem-per-cpu=3600 --account=def-someuser
```

When the job starts, it'll return a prompt on the assigned compute node.

2. On the compute node inside the job **start the ParaView server** using a special version of ParaView with EGL support

```
$ module load paraview/5.13.1
$ unset DISPLAY    # so that PV does not attempt to use X11 rendering context
$ pvserver        # --egl-device-index=0 not needed: first available GPU
                  #      is #0 inside the job
```

For multiple GPUs can use

```
$ nvidia-smi -L    # will return 0, 1, ...
```

The `pvserver` command will return something like

```
Waiting for client...
Connection URL: cs://gra288:11111
Accepting connection(s): gra288:11111
```

Interactive client-server rendering on a cluster's GPU (cont.)

3. On your desktop **set up ssh forwarding** to the ParaView server port:

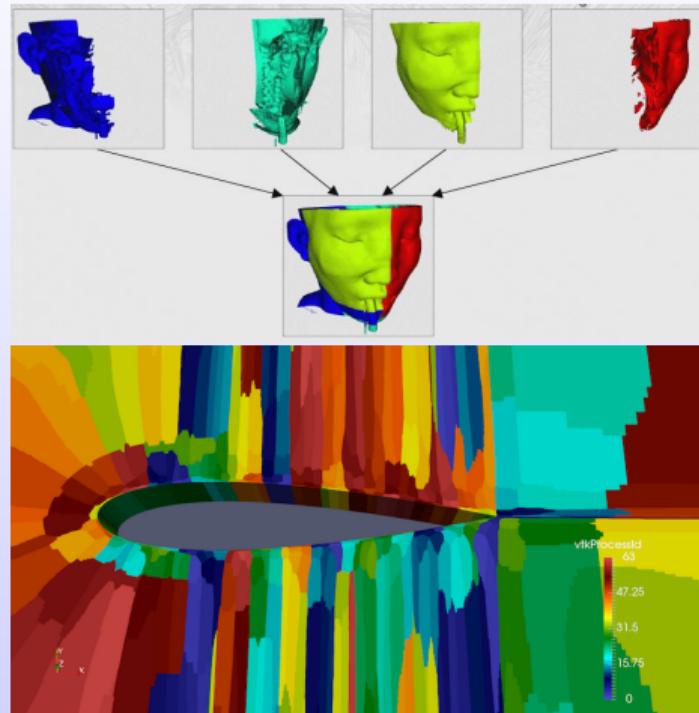
```
$ ssh username@graham.alliancecan.ca -L 11111:gra288:11111
```

4. On your desktop **start ParaView 5.13.x** and **edit its connection properties** under *File - Connect - Add Server* (name = Graham, server type = Client/Server, host = localhost, port = 11111), click *Configure* → *Manual* → *Save*, then select the server from the list and click on *Connect*

-
- ParaView's client and server must have matching major versions (5.13.x)

Data partitioning in parallel ParaView

- If loading unpartitioned data \Rightarrow dynamic load balancing is handled automatically for structured data:
 - structured points
 - rectilinear grid
 - structured grid
- Unpartitioned unstructured data will usually be read in serial, then must be passed through D3 (Distributed Data Decomposition) filter for **dynamic load balancing**:
 - particles/unstructured points
 - polygonal data
 - unstructured grid
- Some unstructured file formats can be read in parallel, e.g. the OpenFOAM reader will automatically read its unstructured data in parallel, distributing it among all available CPU cores
- After passing your unstructured data through D3, you can **save it as parallel PVTU file** \Rightarrow you'll get a **statically distributed dataset** that you can load next time with the same number of CPU cores
- Reading time \Rightarrow you can usually tell if your dataset is being read in serial or in parallel
- Look for `vtkProcessID` variable



Data partitioning in parallel ParaView (cont.)

If you have a large (many GBs) .vtu file:

1. Read your serial .vtu file into parallel ParaView on 16 cores - **slow**
 - and hope that it does not run out of memory on the reading core!
 - at this point the dataset is sitting in memory on one core
 - example: serial .vtu file at 9.1GB \Rightarrow 1'49" reading time
 2. Apply D3 filter to distribute the dataset - **slowish** (memory + MPI)
 3. **File** \rightarrow **Save data** as .pvtu with lz4 level-6 (fast) compression - **fast**
 - \Rightarrow 16 files + 1 header file
 - now you have a statically decomposed dataset
 4. Restart parallel ParaView on 16 cores, read .pvtu from scratch into - **fast!**
 - at this point the dataset is distributed across all 16 cores
 - example: same (but now decomposed) .pvtu dataset at 5.1GB (fast compression) \Rightarrow 11" reading time
- The same I/O speeds logic applies to .vti \rightarrow .pvti (but there is no need for D3)

Exercise: parallel rendering of partitioned data

This is an extremely concise step-by-step guide for the turbine dataset:

1. Submit an interactive job

```
salloc --time=0:60:0 --ntasks=16 --mem-per-cpu=3600
```

2. Start client-server ParaView session on 16 cores

3. Load all .vtm files (all 10 timesteps)

4. Apply Merge Blocks, output type = Unstructured Grid

5. Apply Cell Data to Point Data (so that you could use Contour)

6. Apply D3

7. Save data as decomposed.pvtu, write all timesteps as series, fast compression

8. Restart client-server ParaView session on 16 cores

9. Load all decomposed.pvtu files

10. Create visualization interactively

11. Save animation as 1000×800 PNG files ↗ this step should take $\sim 1\text{min}$ of processing time

12. Merge them into a movie with ffmpeg

Remote rendering summary: some orthogonal decisions

(1) interactive vs. batch

- interactive client-server for a quick look, exploration or debugging
 - another option is to download a scaled-down version of your dataset, debug a script locally on your laptop, and then run it as a batch job on the original full-resolution dataset on the cluster
- batch really preferred for production jobs and producing animations

(2) CPU vs. GPU

- in general, no single answer which one is better
 - you can throw many CPUs at your rendering job
 - modern software rendering libraries such as OSPRay (Intel's ray tracing) and OpenSWR (Intel's rasterizer) can be very fast, depending on your visualization
- might have to resort to software rendering if no GPUs are available (e.g., all are taken by GP-GPU jobs)
- for initial exploration, I would use the dataset size (GBs) to figure out the best number of CPU cores, and adjust from there

SUMMARY

Further resources

- ParaView Discourse

<https://discourse.paraview.org>

- Self-directed ParaView tutorial

<https://docs.paraview.org/en/latest/Tutorials/SelfDirectedTutorial/index.html>

- ParaView User's Guide

<https://docs.paraview.org/en/latest/UsersGuide/index.html>

- ParaView F.A.Q.

<http://www.itk.org/Wiki/ParaView:FAQ>

- VTK wiki with webinars, tutorials, etc. <http://www.vtk.org/Wiki/VTK>

- VTK for C++/Python/Java/C#/JavaScript code examples

<https://kitware.github.io/vtk-examples>

- VTK file formats (3rd-party intro)

<http://www.earthmodels.org/software/vtk-and-paraview/vtk-file-formats>

Our visualization webinars

- ~3-4 visualization webinars per academic year
 - ⌚ subscribe to our mailing list <https://training.westdri.ca/contact>
 - keep an eye on our website <https://training.westdri.ca/blog>
 - ~50 mins + questions, usually on **fairly specific or advanced** topics
- Many past webinars are available with slides and screencasts at <https://bit.ly/vispages>
 - “Launching 2023 Visualize This contest” (showing the Programmable Filter in use)
 - “Image-based approach to large-scale visualization” (Cinema) ● “In-situ visualization with ParaView Catalyst2”
 - “Highlights from the 2021 IEEE SciVis Contest” ● “Text analysis in 3D”
 - “Remote visualization on Compute Canada clusters” ● “Data visualization in Julia with the Makie ecosystem”
 - “Scientific visualization on NVIDIA GPUs”
 - “Workflows with Programmable Filter / Source in ParaView”
 - “The Topology ToolKit (TTK)” ● “Command-line image processing with ImageMagick”
 - “Web-based 3D scientific visualization” (ParaViewWeb, vtk.js, ParaView Glance)
 - “Photorealistic rendering with ParaView and OSPRay”
 - “Batch visualization on Compute Canada clusters”
 - “Molecular visualization with VMD” ● “Intermediate VMD topics: trajectories, movies, scripting”
 - “Using YT for analysis and visualization of volumetric data” (part 1) ● “Working with data objects in YT” (part 2)
 - “Scientific visualization with Plotly”
 - “Novel visualization techniques from 2017 VISUALIZE THIS competition”
 - “Camera animation in ParaView and VisIt”
 - “3D visualization on new Compute Canada systems”
 - “Using ParaViewWeb for 3D visualization and data analysis in a web browser”
 - “Visualization support in WestGrid / Compute Canada”
 - “Scripting and other advanced topics in VisIt visualization”
 - “CPU-based rendering with OSPRay”
 - “3D graphs with NetworkX, VTK, and ParaView” ● “Graph visualization with Gephi”
- We are always looking for topic suggestions!

Documentation and getting help

- Visualization in the Alliance <https://ccvis.netlify.app>
(online gallery)
- Official documentation <https://docs.alliancecan.ca/wiki/Visualization>
- Western Canada research computing visualization resources <https://bit.ly/vispages>
(webinar archive)
- Email **support@tech.alliancecan.ca** and mention “*visualization*” in the subject line (goes to our ticketing system)
- Email me **alex.razoumov@westdri.ca**
- ParaView documentation
 - official documentation <https://docs.paraview.org/en/latest>
 - wiki <http://www.paraview.org/Wiki/ParaView>
 - Python batch scripting <http://bit.ly/2wF5v0B>
 - VTK tutorials <http://www.itk.org/Wiki/VTK/Tutorials>